

Synthetic BOAMP Benchmark for Renewal-Linkage Evaluation

Latent-Truth Generator, Conditional Fidelity Validation, Temporal and Candidate-Environment Readiness

Gigalis Predictive-Modeling Internship — Current Independent Study

Repository state: July 2026; benchmark version `v0_3_temporal_candidate_revision`

Abstract

This report documents the synthetic BOAMP benchmark now used to evaluate procurement-renewal linkage algorithms when legal ground-truth renewal labels are unavailable in the real BOAMP corpus. The benchmark is not a substitute for the original dataset. Its purpose is narrower and more defensible: it creates a controlled world with hidden recurrence truth, corrupts the observable fields to resemble real BOAMP notices, and then validates whether the resulting observed synthetic data is realistic enough for testing candidate generation and linkage behavior.

The current benchmark version, `v0_3_temporal_candidate_revision`, passes the readiness gates required before running linkage algorithms. Conditional field fidelity remains passing after the v0.2 revision; temporal fidelity passes at the publication-year, schema-year, notice-type-year, and 12/24-month follow-up levels; the Layer-1 production candidate environment now passes after the v0.3 revision; and structural truth checks confirm that hidden labels are internally coherent and not leaked into observed notices. The candidate environment improved from a v0.2 zero-candidate rate of 90.4% to 63.0%, compared with 60.9% in the real Layer-1 source scope. The final status is **READY_FOR_LINKAGE**, with a documented caveat that 60-month follow-up runway remains weaker than the real corpus and that synthetic-benchmark evidence cannot replace inaccessible real truth.

Contents

1	Technical Summary	2
2	Problem and Evidence Standard	3
2.1	The real-data limitation	3
2.2	Definitions of important terms	3
3	Benchmark Data Scope and Artifacts	4
3.1	Current benchmark version	4
3.2	Files to use for linkage evaluation	4
4	End-to-End Generator Pipeline	5
4.1	Flowchart	5
4.2	Pipeline stages	5
4.3	Concrete generation procedure	6
5	Latent-Truth Model	7
5.1	Buyer population	7
5.2	Procurement needs	7
5.3	Recurrence propensity	7
5.4	Cycle duration and successor gaps	8
5.5	v0.3 digital-scope recurrence timing	8

6	Observed-Corruption Model	8
6.1	Why corruption is separate from truth	8
6.2	Latent quality class and co-occurring errors	8
6.3	Weighted categorical corruption	9
6.4	Conditional field observation	9
6.5	Text mechanisms	9
7	Validation Methodology	10
7.1	Weighted mean absolute error	10
7.2	Total variation distance	10
7.3	Candidate-count percentiles	10
7.4	Evaluation workflow	10
8	Revision History	11
8.1	v0.1 provisional benchmark	11
8.2	v0.2 conditional revision	11
8.3	v0.3 temporal/candidate revision	11
9	Final Results	12
9.1	Freeze-gate status	12
9.2	Conditional fidelity	12
9.3	Temporal validation	12
9.4	Candidate environment	12
9.5	Structural truth	13
10	Why the Benchmark Is Ready for Linkage Algorithms	13
11	Limitations and Caveats	14
12	Reproducibility and Source Inventory	14
13	Recommended Next Step	15

1 Technical Summary

The practical problem is that BOAMP notices do not contain a reliable field saying that one call for tenders is the renewal or recurrence of a previous procurement need. The production linkage algorithm must infer that relation from observable signals such as buyer identity, CPV code, object text, and timing. Because the true relation is unavailable in the real dataset, algorithm evaluation needs a controlled benchmark where the true relation is known but hidden from the algorithm.

The implemented benchmark follows a *latent truth plus observed corruption* design. First, a clean synthetic world is generated: buyers, establishments, procurement needs, contract cycles, notices, and true successor relations. Second, observable BOAMP fields are corrupted or omitted according to calibrated and scenario-documented mechanisms. Third, the observed synthetic notices are validated against the real corpus, while the hidden truth table is kept separate for later algorithm evaluation.

Current decision. The v0.3 central-provisional `observed_notices.parquet` file is ready to be used as the observed input to the linkage algorithms. The corresponding hidden labels are in `true_relations.parquet`. The linkage algorithms must not read the truth file during candidate generation, scoring, calibration, or threshold selection.

The final benchmark is credible enough for linkage testing for four reasons.

1. **Conditional fidelity passes.** The benchmark reproduces key observable conditional rates: SIRET presence, duration presence, CPV missingness, and generic or repeated text.
2. **Temporal fidelity passes.** Publication-year composition, schema-by-year composition, notice-type-by-year composition, and 12/24-month follow-up runway match the real corpus within the defined acceptance gates.
3. **Candidate environment passes.** The actual Layer-1 candidate generator sees a synthetic source/candidate environment close enough to the real production source scope.
4. **Structural truth passes.** Every cycle has exactly one outgoing truth relation, observed notices contain no truth columns, synthetic identifiers are valid and non-leaking, and a meaningful broad true-gap tail remains.

2 Problem and Evidence Standard

2.1 The real-data limitation

BOAMP is the *Bulletin Officiel des Annonces de Marches Publics*, the French public-procurement gazette. It publishes standalone notices such as *appel d’offre* notices (calls for tenders) and *attribution* notices (award notices). A later notice may correspond to a recurrence of an earlier procurement need, but the corpus does not provide a complete ground-truth relation table.

This creates an evaluation problem. If an algorithm links notice i to notice j , the real dataset generally cannot say with certainty whether that link is correct. Manual review can help, but it is limited and expensive. The synthetic benchmark therefore provides a controlled evaluation environment:

known latent truth $\rightarrow$ realistic observed corruption $\rightarrow$ blind linkage evaluation.

The benchmark is not intended to prove real-world renewal rates. It is intended to test whether a linkage algorithm can recover known recurrence structure under observation patterns similar to the real BOAMP data.

2.2 Definitions of important terms

Table 1: Technical terms used in this report.

Term	Meaning
BOAMP	French public-procurement notice corpus. In this project, it is the real source used to calibrate observable field behavior.
CPV	Common Procurement Vocabulary. It is a European classification code for procurement subject matter. Digital/ICT scope uses divisions 32, 35, 48, and 72.
SIREN	French nine-digit organization identifier. It identifies a legal unit.
SIRET	French fourteen-digit establishment identifier. It begins with the corresponding SIREN and identifies a specific establishment.
Latent truth	The clean hidden synthetic world: true buyers, establishments, needs, cycles, durations, and successor relations.
Observed corruption	The process that transforms clean truth into observed fields with missingness, invalid identifiers, rounded duration, generic text, or name variation.

Term	Meaning
Candidate environment	The set of possible future notices produced by the production candidate generator for each source notice. It is evaluated before scoring.
Zero-candidate rate	The share of eligible source notices for which the candidate generator finds no candidate.
Percentile / p90	The 90th percentile. If p90 candidate count is 6, then 90% of sources have at most 6 candidates.
WMAE	Weighted mean absolute error. It measures average absolute real-versus-synthetic error while weighting each cell by the real cell size.
TVD	Total variation distance. It measures the distance between two discrete distributions.
Structural truth	Internal validity of hidden labels: one outgoing relation per cycle, no duplicated successor target where forbidden, no truth leakage into observed data.
Truth leakage	A failure mode where fields such as <code>cycle_id_true</code> or <code>need_id_true</code> appear in the observed data used by the algorithm.

3 Benchmark Data Scope and Artifacts

3.1 Current benchmark version

The active benchmark version is `v0_3_temporal_candidate_revision`. It keeps the v0.2 conditional-field improvements and adds a temporal/candidate-environment revision targeted at the actual Layer-1 source scope. The benchmark was generated with fixed seeds:

latent world seed = 20260721, corruption seed = 20260722.

Table 2: Generated v0.3 central-provisional row counts. Counts come from `generation_metadata.json`.

Table	Rows
Latent buyers	2,000
Latent establishments	2,174
Latent procurement needs	4,800
Latent cycles	6,866
True relations	6,866
Notice-family membership	9,471
Clean notices	9,471
Observed notices	9,471
Corruption log	32,344

3.2 Files to use for linkage evaluation

Table 3: Primary v0.3 artifacts.

Artifact	Role
<code>observed notices</code>	The observed synthetic notice table, stored as <code>observed_notices.parquet</code> . This is the input that linkage algorithms may read.
<code>true relations</code>	Hidden evaluation labels, stored as <code>true_relations.parquet</code> . This must be used only after algorithm outputs are produced.
<code>generation metadata</code>	Generator version, seeds, config hashes, row counts, selected v0.3 candidate-revision parameters, and validation status.
<code>corruption log</code>	Audit trail showing how clean fields were corrupted into observed fields. It is for generator validation, not linkage input.
<code>selected params</code>	Parameter-provenance CSV for the v0.3 temporal/candidate revision.

4 End-to-End Generator Pipeline

4.1 Flowchart

4.2 Pipeline stages

The generator is split into two independent random streams: one for the latent clean world and one for observation/corruption. This separation lets a future experiment hold the true world fixed while changing observation corruption, or hold corruption fixed while changing latent truth assumptions.

Table 4: Generator stages and their interpretation.

Stage	Description
Latent buyers	Generates synthetic public buyers with department, buyer type, synthetic organization name, SIREN, activity rate, activity tier, active window, alias propensity, and identifier-quality propensity. Buyer activity is heavy-tailed, so concentration is emergent rather than assigned directly.
Latent establishments	Generates valid SIRET values under the buyer’s SIREN hierarchy. This permits establishment-level identity errors without using real identifiers.
Latent needs	Generates procurement needs under each buyer: CPV code, broad segment, base text concepts, vocabulary, duration profile, and recurrence propensity.
Latent cycles	Generates one or more contract cycles for each need. Each cycle has a start date, true duration, expected end, and possibly a successor cycle.
True relations	Derives <code>NEXT_CYCLE</code> or <code>NO_SUCCESSOR</code> for every cycle from the latent cycle table. This is the hidden evaluation truth.

Stage	Description
Clean notices	Generates an <code>APPEL_OFFRE</code> call notice for every cycle and an <code>ATTRIBUTION</code> award notice with calibrated probability.
Observed corruption	Corrupts identifiers, buyer names, CPV, duration, object text, and linked-call visibility. This creates observed notices shaped like real prepared BOAMP data.
Validation	Compares observed synthetic data against real data on field availability, temporal structure, candidate environment, and structural truth.

4.3 Concrete generation procedure

The implemented generation process can be read as the following reproducible algorithm. This is the logic used to create the final v0.3 central-provisional dataset; row counts are emergent outputs, not manually forced row-by-row.

1. **Load fixed inputs.** The run loads the benchmark defaults, the central-provisional scenario, empirical calibration tables, and the prepared real BOAMP corpus used to estimate observable rates. The v0.3 run uses fixed seeds 20260721 for the latent world and 20260722 for corruption.
2. **Generate buyers.** The generator samples 2,000 synthetic buyers with department, buyer type, synthetic name, SIREN, active date window, activity rate, identifier-quality propensity, and alias propensity. No real buyer name or real identifier is copied.
3. **Generate establishments.** For each buyer, one or more valid SIRET establishments are created under that buyer’s SIREN. This produces a realistic SIREN–SIRET hierarchy before any observed identifier corruption is applied.
4. **Generate procurement needs.** For each buyer, the number of needs is sampled from the activity-tier Poisson model. Each need receives a synthetic CPV code, a CPV-conditioned text vocabulary, a duration profile, and a recurrence propensity.
5. **Generate cycles.** Each need starts with one procurement cycle. Additional cycles are sampled from the recurrence propensity until the maximum cycle count, failed successor draw, or observation-window boundary stops the chain. The cycle table contains true start dates, true durations, and true expected end dates.
6. **Derive hidden truth.** The truth table is derived from cycles only. If a cycle is followed by another cycle of the same need, the relation is `NEXT_CYCLE`; otherwise it is `NO_SUCCESSOR`. Thus every cycle has exactly one outgoing truth relation.
7. **Generate clean notices.** Every cycle creates one call-for-tenders notice. An award notice is added with the calibrated award probability. These notices still contain clean truth fields and are not given to the linkage algorithm.
8. **Corrupt observed fields.** Identifier, buyer-name, CPV, duration, text, and linked-call visibility mechanisms transform clean notices into observed notices. The observed output uses real-pipeline-like columns and intentionally contains no truth columns.
9. **Apply v0.3 candidate revision.** For the digital CPV divisions 32, 35, 48, 72, selected v0.3 settings increase the chance that future same-buyer candidates exist in the production candidate window while preserving a broad true-gap tail and relation integrity.

10. **Write versioned outputs.** The final files are written under `v0_3_temporal_candidate_revision`; v0.2 outputs are not overwritten, so before/after comparisons remain reproducible.

This procedure explains why the benchmark contains both `observed_notices.parquet` and `true_relations.parquet`. The first file is what the algorithm sees. The second file is hidden truth used only for evaluation after predictions are produced.

5 Latent-Truth Model

5.1 Buyer population

Buyer activity is sampled from a heavy-tailed Pareto distribution. If A_b is the raw activity draw for buyer b , then

$$A_b \sim \text{Pareto}(1.0) + 0.15, \quad r_b = \frac{A_b}{\sum_{b'} A_{b'}}.$$

The normalized value r_b is an activity rate. It is not directly used as a true renewal label; it shapes how many needs and notices a buyer tends to have. The intent is to reproduce the real corpus’s buyer concentration pattern without copying real buyer names or identifiers.

Each buyer also receives:

$$q_b \sim \text{Beta}(5, 2), \quad a_b \sim \text{Beta}(2, 5),$$

where q_b is an identifier-quality propensity and a_b is an alias-propensity value used later by the corruption layer.

5.2 Procurement needs

For each buyer b , the number of latent procurement needs is

$$N_b = \max\{1, \text{Poisson}(\lambda_{\text{tier}(b)})\}. \quad (1)$$

The implemented tier parameters are

$$\lambda_{21+} = 8.0, \quad \lambda_{6-20} = 4.0, \quad \lambda_{2-5} = 1.8, \quad \lambda_1 = 1.2.$$

Each need receives a CPV division sampled from the empirical real-corpus division distribution, an eight-digit synthetic CPV-looking code, a broad text vocabulary, a duration profile, and a recurrence propensity.

The per-need duration profile is

$$D_{\text{profile}} \sim \text{LogNormal}(\log 12, 0.7), \quad (2)$$

clipped to the project duration range where needed. This creates positive, right-skewed contract-duration tendencies rather than one fixed duration for all needs.

5.3 Recurrence propensity

The central provisional scenario uses a base recurrence propensity $p_0 = 0.40$. This is a scenario assumption, not a real-world estimate. Each need receives multiplicative jitter

$$J = 0.7 + 0.6B, \quad B \sim \text{Beta}(4, 4), \quad (3)$$

so the implemented recurrence probability is

$$p_{\text{need}} = \text{clip}(p_0 J, 0, 0.98). \quad (4)$$

The cap at 0.98 avoids deterministic recurrence while allowing highly recurrent needs.

5.4 Cycle duration and successor gaps

For each realized cycle, the true duration is drawn as

$$D_{\text{cycle}} \sim \text{clip}_{[1,120]} \{ \mathcal{N}(D_{\text{profile}}, (0.25D_{\text{profile}})^2) \}. \quad (5)$$

The expected true end date is the true start date plus D_{cycle} months, approximated by 30.44 days per month in implementation.

Before v0.3’s scoped candidate revision, broad successor gaps used the central distribution

$$G_{\text{broad}} \sim \text{clip}_{[1,\infty)} \{ \mathcal{N}(18, 9^2) \}. \quad (6)$$

This distribution is deliberately not centered inside the production six-month candidate window. The reason is important: latent truth should not be tuned only to make the algorithm easy. A meaningful share of true successors must remain outside the six-month window so the benchmark can test limitations in candidate generation.

5.5 v0.3 digital-scope recurrence timing

The v0.3 revision added an optional scoped mechanism for the digital CPV divisions used by the production source filter:

$$\mathcal{D}_{\text{digital}} = \{32, 35, 48, 72\}.$$

For scoped digital needs, recurrence propensity is multiplied by 1.3 and capped at 0.98. When a successor gap is drawn, the selected v0.3 parameter set uses a near-window gap with probability 0.75:

$$G_{\text{near}} \sim \text{clip}_{[0.1,6]} \{ \mathcal{N}(2, 2^2) \}. \quad (7)$$

Otherwise the broad gap G_{broad} is used. This creates more future same-buyer candidates inside the production candidate window without forcing candidate counts directly onto rows.

The v0.3 revision also permits limited same-buyer hard-negative chain alignment at rate 0.10. This shifts whole distinct-need chains from the same buyer near a source’s expected end date. Because entire chains are shifted together, within-need cycle order and truth relations remain valid. The aligned target is never the same `need_id_true`, so these are plausible hard negatives rather than mislabeled true successors.

6 Observed-Corruption Model

6.1 Why corruption is separate from truth

The report’s central methodological distinction is:

$$\text{latent true value} \neq \text{observed BOAMP field}.$$

For example, every latent cycle has a true duration, but the observed `declared_duration_months` field may be missing, rounded, converted incorrectly, or imputed later by the compatibility adapter. This separation prevents a common modeling mistake: changing the true contract duration just to fix observed duration missingness.

6.2 Latent quality class and co-occurring errors

Each notice is assigned a latent quality class $Q_i \in \{\text{HIGH}, \text{MEDIUM}, \text{LOW}\}$. The scenario class mix is adjusted by the buyer’s identifier-quality propensity. Field corruption rates are then scaled by quality-class multipliers:

$$m(Q_i) = \begin{cases} 0.35, & Q_i = \text{HIGH}, \\ 1.00, & Q_i = \text{MEDIUM}, \\ 2.20, & Q_i = \text{LOW}. \end{cases}$$

Duration uses a flatter multiplier set $\{0.89, 1.00, 1.15\}$, because real duration presence behaves more like a schema/template phenomenon than a general record-quality phenomenon. This design makes field errors co-occur on lower-quality records instead of sampling every error independently.

6.3 Weighted categorical corruption

Several corruption mechanisms use the same pattern. If a field has base error rates r_k for outcomes k , the quality-scaled rates are

$$w_k = \min(0.98, r_k m_i),$$

renormalized if their sum exceeds one. The remaining probability is assigned to the non-corrupted outcome:

$$P(\text{OK}) = 1 - \sum_k w_k.$$

This is used for mechanisms such as old identifier corruption and CPV corruption.

6.4 Conditional field observation

The v0.2 revision replaced mainly global field mechanisms with conditional observation targets estimated from observable real data. For a cell g , such as schema family by year by notice type, the smoothed target rate is

$$\hat{p}_g = \frac{x_g + \alpha \hat{p}_{\text{global}}}{n_g + \alpha}, \quad \alpha = 25, \quad (8)$$

where x_g is the number of observed successes in the real cell, n_g is the real cell size, and $\hat{p}_{\text{global}}$ is the real global rate. This is a beta-binomial style shrinkage estimator. It prevents sparse cells from being copied too noisily while preserving large-cell conditional structure.

The implemented conditional targets are:

$$P(\text{SIRET present} \mid \text{schema, year, notice type}),$$

$$P(\text{duration present} \mid \text{schema, year, notice type}),$$

and

$$P(\text{generic repeated text} \mid \text{notice type, buyer activity tier}).$$

CPV missingness is modeled by schema and notice type and remained passing without v0.3 retuning.

6.5 Text mechanisms

The generator separates several text concepts that should not be collapsed:

1. base procurement text generated from CPV-division vocabulary;
2. small same-cycle variation between call and award notices;
3. next-cycle lexical drift for recurrent needs;
4. exact cross-buyer templates;
5. near-boilerplate text with buyer-, CPV-, date-, or department-specific tokens;
6. weak generic procurement text.

This distinction matters because exact duplicate text across unrelated buyers is a different challenge from realistic repeated administrative language within the same buyer or same procurement family.

7 Validation Methodology

7.1 Weighted mean absolute error

For conditional rates, the main fidelity metric is weighted mean absolute error:

$$\text{WMAE} = \frac{\sum_g n_g^{\text{real}} |\hat{p}_g^{\text{syn}} - \hat{p}_g^{\text{real}}|}{\sum_g n_g^{\text{real}}}. \quad (9)$$

The unit is percentage points when rates are percentages. Weighting by n_g^{real} is essential because small real cells can have noisy rates. A rare cell with 20 records should not count as much as a common cell with 10,000 records.

7.2 Total variation distance

For distribution comparisons such as publication-year share, the report uses total variation distance:

$$\text{TVD}(p, q) = \frac{1}{2} \sum_k |p_k - q_k|. \quad (10)$$

A value of zero means the two discrete distributions match exactly. Larger values mean more mass would need to be moved across categories to make them match.

7.3 Candidate-count percentiles

If C_i is the candidate count for source notice i , then the empirical q -th percentile is

$$P_q(C) = \inf\{c : \hat{F}_C(c) \geq q\}. \quad (11)$$

In plain language, p90 is the value below which 90% of sources fall. For example, real p90 candidate count = 6 means that 90% of real Layer-1 source notices have at most six candidates.

The zero-candidate rate is

$$z = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(C_i = 0). \quad (12)$$

This is a critical candidate-environment metric because a source with no candidate cannot be recovered by any downstream scoring model.

7.4 Evaluation workflow

The validation is performed in three layers before the benchmark is declared ready for linkage.

1. **Conditional field validation.** The notebook compares real and synthetic conditional probabilities for SIRET presence, duration presence, CPV missingness, and repeated/generic text. It reports cell-level tables, weighted and unweighted errors, signed bias, p90 error, maximum error, and pass/fail status.
2. **Temporal validation.** The notebook compares the calendar structure of real and synthetic notices: publication-year share, publication-month share, schema by year, notice type by year, and follow-up runway at 12, 24, 36, and 60 months.
3. **Candidate-environment validation.** The observed synthetic notices are adapted into the same shape expected by the production Layer-1 candidate generator. The actual candidate generator is then run, and the resulting source counts, zero-candidate rate, candidate-count percentiles, cap-reached rate, and buyer-key-specific zero-candidate rates are compared with the real Layer-1 source scope.

4. **Structural-truth validation.** The validation checks that truth labels are coherent and hidden: one outgoing relation per cycle, no truth columns in observed notices, no identifier leakage, metadata validation status **PASS**, and broad true-gap tail above the required threshold.

The key point is that the candidate environment is evaluated through the same production candidate-generation logic used later by the linkage algorithm. Candidate counts are fidelity outcomes, not direct generator assignments. This prevents the benchmark from being tuned to make a particular scorer look good.

Table 5: Evaluation layers and their evidence outputs.

Evaluation layer	Main question	Main saved evidence
Conditional fidelity	Do observed fields appear with realistic conditional rates?	Conditional fidelity summary and all-cell tables.
Temporal fidelity	Does the synthetic corpus have realistic calendar and follow-up structure?	Temporal summary and follow-up runway tables.
Candidate environment	Does the production candidate generator see a realistic source/candidate universe?	Candidate environment summary and gate tables.
Structural truth	Is hidden truth coherent and absent from observed notices?	Temporal truth sanity checks and metadata.

8 Revision History

8.1 v0.1 provisional benchmark

The v0.1 benchmark established the initial latent-truth-and-corruption architecture. It generated synthetic buyers, needs, cycles, notices, true relations, and observed corruptions. It also created initial global fidelity checks. The main limitation was that several fields matched marginal rates more than conditional structure.

8.2 v0.2 conditional revision

The v0.2 revision targeted conditional field fidelity. It replaced or refined mechanisms for SIRET presence, duration presence, CPV missingness, and repeated/generic text. The most important methodological improvement was moving from global field rates to smoothed conditional observation models based on observable real-data cells.

The v0.2 result passed conditional fidelity but failed the production candidate environment. In the actual Layer-1 source scope, the synthetic zero-candidate rate was 90.4%, far above the real 60.9%, and p90 candidates/source was 0 instead of the real value 6. This meant the benchmark was not yet ready for linkage evaluation, even though the field-level validation had improved.

8.3 v0.3 temporal/candidate revision

The v0.3 revision preserved v0.2 conditional fidelity and changed the latent digital recurrence/candidate environment. It did not assign candidate counts directly. It did not use linkage scores, accepted links, precision, recall, score calibration, or thresholds as generator inputs.

The bounded pilot sweep tested fixed ordered parameter sets for recurrence multiplier, near-

window share, and hard-negative alignment rate. The first passing setting was:

$$\begin{aligned} \text{recurrence multiplier} &= 1.3, \\ \text{near-window share} &= 0.75, \\ \text{hard-negative alignment rate} &= 0.10. \end{aligned}$$

Additional selected settings were same-buyer clustering for scoped digital needs, scoped buyer affinity share 0.07, high scoped-need probability 0.70, stable scoped establishment, and stable scoped name fallback.

9 Final Results

9.1 Freeze-gate status

Table 6: Final v0.3 freeze gate.

Dimension	Status	Interpretation
Conditional fidelity	PASS	v0.3 remains under thresholds with no material regression versus v0.2.
Temporal validation	PASS	Calendar and 12/24-month follow-up gates pass.
Candidate environment	PASS	Production Layer-1 candidate generator passes source-scope gates.
Structural truth	PASS	Metadata, truth leakage, and broad-gap-tail checks pass.
Linkage readiness	READY FOR LINK-AGE	Linkage algorithms may now be run on observed v0.3 notices.

9.2 Conditional fidelity

Table 7: Conditional fidelity summary. WMAE values are in percentage points.

Target	v0.2 WMAE	v0.3 WMAE	Regression vs v0.2	Status
CPV missingness	2.57	2.09	-0.48	PASS
Duration presence	0.51	0.62	+0.11	PASS
Generic repeated text	1.40	0.80	-0.60	PASS
SIRET presence	1.88	1.60	-0.28	PASS

9.3 Temporal validation

9.4 Candidate environment

The key improvement is the zero-candidate rate. In v0.2, most synthetic sources had no possible candidates, so downstream scoring algorithms would mostly be tested on an unrealistically empty environment. In v0.3, the zero-candidate rate is within 2.16 percentage points of the real Layer-1 source scope. The upper candidate-count tail remains lighter than the real corpus, but it passes the predefined p75, p90, and p95 gates.

Table 8: Temporal summary metrics for v0.3.

Dimension	Metric	Value	Status
Publication-year share	TVD	0.151	PASS
Schema by year	WMAE	0.34 pp	PASS
Notice type by year	WMAE	3.80 pp	PASS
Publication-month share	TVD	0.178	PASS

Table 9: Follow-up runway comparison. The 60-month row remains a documented caveat.

Horizon	Real evaluable	Synthetic evaluable	Absolute error	Status
12 months	91.68%	90.82%	0.85 pp	PASS
24 months	82.52%	81.12%	1.40 pp	PASS
36 months	73.23%	70.19%	3.03 pp	PASS
60 months	55.02%	44.49%	10.52 pp	Caveat

Table 10: Layer-1 algorithm source-scope candidate environment.

Scope	Sources	Zero-cand.	Mean	p75	p90	p95
Real	3,159	60.9%	1.94	2	6	10
v0.2 synthetic	354	90.4%	0.15	0	0	1
v0.3 synthetic	395	63.0%	0.64	1	2	3

Table 11: Structural-truth checks for v0.3.

Metric	Value	Status
Cycles	6,866	PASS
Successor relation rate	30.1%	Diagnostic
Share of true gaps over 6 months	81.2%	PASS
Observed truth columns	0	PASS
Metadata validation status	PASS	PASS

9.5 Structural truth

The broad true-gap tail is especially important. A share of 81.2% of true successor gaps remains above six months, well above the minimum gate of 45%. This means the revision did not collapse all truth into the production candidate window merely to make the candidate generator look better.

10 Why the Benchmark Is Ready for Linkage Algorithms

The benchmark is ready because the validation evidence now matches the task the linkage algorithms will face. Earlier versions could match some global or conditional field statistics but still fail to produce a realistic candidate environment. That is not sufficient: a linkage algorithm first needs plausible candidates before scoring can matter.

The v0.3 benchmark passes the hard readiness gates:

- observed field availability and text reuse are close to real conditional patterns;
- temporal notice distribution is close enough for the planned evaluation stage;
- the actual Layer-1 candidate generator produces a realistic enough zero-candidate rate and candidate-count distribution;
- hidden truth exists, is internally coherent, and is not present in the observed notices;
- v0.3 did not use linkage thresholds or accepted links as tuning targets.

Therefore, the next valid step is to run the linkage algorithms on `observed_notices.parquet`, produce predicted links, and only then compare those predictions against `true_relations.parquet`. Linkage-score calibration, threshold optimization, and precision/recall interpretation belong to that next stage, not to the benchmark-generation stage.

11 Limitations and Caveats

The benchmark has passed the current readiness gate, but several limitations remain.

1. **Synthetic evidence is not real legal truth.** It can test whether an algorithm behaves sensibly in a controlled realistic environment, but it cannot prove real-world renewal accuracy.
2. **Some latent quantities are scenario assumptions.** Recurrence prevalence, broad successor-gap distribution, and text drift are documented assumptions where real truth is unavailable.
3. **The 60-month follow-up runway is weaker.** The 12- and 24-month gates pass, but the 60-month comparison remains a caveat.
4. **Candidate tail is lighter than real.** v0.3 passes p75/p90/p95 gates, but real p90 and p95 counts remain higher than synthetic values.
5. **No score calibration has been done yet.** The benchmark is ready for running linkage algorithms, not for claiming final linkage performance.

12 Reproducibility and Source Inventory

The main notebook for the final validation is `notebooks/10_synthetic_temporal_candidate_revision_validation.ipynb`. The concise generated report is `reports/generated/synthetic_benchmark/v0_3_temporal_candidate_revision_report.md`. The final machine-readable readiness gate is `reports/tables/synthetic_benchmark/v0_3_temporal_candidate_revision/temporal_candidate_freeze_gate.csv`. The selected parameter provenance is `reports/tables/synthetic_benchmark/v0_3_temporal_candidate_revision/candidate_revision_selected_parameters.csv`.

The primary implementation modules reviewed for this report are:

- `src/boamp/synthetic/buyers.py`;
- `src/boamp/synthetic/needs.py`;
- `src/boamp/synthetic/cycles.py`;
- `src/boamp/synthetic/notices.py`;
- `src/boamp/synthetic/corruption.py`;

- `src/boamp/synthetic/conditional_observation.py`;
- `scripts/generate_synthetic_benchmark_v0_3_temporal_candidate_revision.py`.

13 Recommended Next Step

The correct next step is to run linkage algorithms against the v0.3 observed synthetic notices, store their predicted links, and evaluate them against the hidden truth relations. The evaluation should report candidate coverage, ranking quality, precision/recall against hidden truth, threshold sensitivity, and failure cases by buyer-key type, CPV division, text pattern, and temporal gap.

Until that evaluation is complete, the defensible claim is:

The synthetic benchmark is credible enough to begin linkage-algorithm evaluation under controlled truth, but it does not replace inaccessible real renewal labels.




Figure 1: End-to-end benchmark pipeline. The blue layer contains hidden truth and must not be used by the linkage algorithm. The orange layer is the observed synthetic data that mimics real BOAMP field availability and corruption.




Figure 2: Conditional-field fidelity after the v0.3 revision. The figure shows that v0.3 preserves the v0.2 conditional fixes. It supports field-fidelity readiness, but it does not by itself prove candidate-environment realism.




Figure 3: Publication-year composition. This visual checks whether synthetic notices are spread across years similarly to the real corpus. It is a calendar-structure check, not a linkage-accuracy check.




Figure 4: Publication-month composition. This supports the seasonal/calendar plausibility of the benchmark.




Figure 5: Follow-up runway comparison. The 12- and 24-month horizons pass; the 60-month horizon remains visibly weaker and should be stated as a limitation in downstream linkage evaluation.




Figure 6: Candidate environment before and after v0.3. This is the most important visual for linkage-readiness: it shows that the actual production candidate generator now sees a realistic enough source/candidate environment.